

Minimum Variance

Problem description

- Network of N nodes. A given node N_i
 - has a fixed upload capacity, C_i
 - has a variable number of connections, K_i
- Network is an $N \times N$ matrix of binary values (i.e. connected or not)
- We have a fitness function, $F(m)$
 - F takes a $N \times N$ binary matrix
 - and returns the variance in download rates

Brute-force solution

- A brute force solution would test every possible state
- For a given row, R_i
 - Number of combinations is equal to the number of subsets
 - Thus $2^m - 1$, where $m = \min(N, C_i/5)$
- For entire matrix, number of combinations close to 2^{n*n}
- Not feasible to calculate directly

So how about we rephrase this as an
optimisation problem?

Optimisation problem

- Attempts to maximise (or minimise) a function, $g(x)$
 - where x is a discrete state of the problem
- Perfect fit for our situation
 - we have well defined state, i.e. the binary matrix
 - and a function to minimise, $F(m)$
- Many techniques to solve or approximate this class of problem
 - Hill climbing
 - Random walks
 - Simulated annealing
 - etc.

Lets look at hill climbing

Hill climbing

- Start with some state
 - doesn't matter which, can be random
- Each cycle, pick a neighbouring state to move to
 - pick neighbouring state with best value of fitness function
- When we can't find a better neighbour, we have found a maximum
 - Unfortunately, it might just be a local maxima
- We can alleviate the problem of local maxima
 - Stochastic hill climbing: pick a random better neighbour
- Still not guaranteed to find an optimal solution
 - But might give us a helpful approximation

I implemented a stochastic hill climber
but it always gets stuck in local minima

We need something more reliable
Perhaps simulated annealing?

But wait, what does our simulation do?

Our simulation

- Represents state as a sparse $N \times N$ matrix
- Each node attempts to maximise its own download rate
 - by moving to neighbouring state
 - and regularly reevaluating metric
 - basically, hill climbing on smaller scale
- Maximising majority of download rates \sim minimising variance
 - Because large peers unable to maximise their download
 - Therefore, we are optimising to minimise variance
- Unlike a hill climber, we allow backtracking
 - never converges to a fixed solution
 - never gets stuck either
- I believe our simulation is approximating minimum variance